

# ADT302 - CONCEPTS IN BIG DATA ANALYTICS

Viju P Poonthottam  
Asst. Professor  
Dept. of AI & DS  
MES CE Kuttippuram

March 24, 2023



Table of contents


SYLLABUS

History

Anatomy of a File Read

## *Module 3 -Hadoop Distributed File System(13 Hrs)*

- History of Hadoop
- Hadoop Ecosystem and Core Components
- HDFS Architecture
- Using HDFS Files ,HDFS Design
- Blocks, Namenodes and Data nodes
- Basic File system Operations
- Hadoop Specific File Types
- Anatomy of a file read
- Anatomy of a file write
- Execution pipeline
- Runtime Coordination and Task Management in MapReduce
- Using MapReduce as a framework for parallel processing
- Road Enrichment Example



## *History*

- The project's creator, Doug Cutting
- kid gave a stuffed yellow elephant. Short, relatively easy to spell and pronounce, meaningless, and not used elsewhere
- GFS
- In 2004, Google published the paper that introduced MapReduce
- Eric Baldeschwieler- new framework, written in C++ modeled with GFS and MapReduce
- Yahoo! hired Doug Cutting,
- Hadoop was made its own top-level project at Apache
- Hadoop was being used by many other companies besides Yahoo!, such as Facebook, and the New York Times etc.



- HDFS - Hadoop Distributed file system
- MapReduce -Is a programming model for processing large data sets

# *HDFS Structure*


- Namenode
- Datanode
- HDFS Client




## *Block Placement*

- Blocks – Default size 64MB
- Replica


# *HDFS*



# *HDFS*




# HDFS




HDFS stores these  
blocks on datanodes

## *HDFS*



HDFS distributes the  
blocks to the DN<sub>s</sub>

# HDFS



The NameNode tracks  
blocks and Datanodes

## *Basic File system Operations*

- `hadoop fs -copyFromLocal input/docs/quangle.txt hdfs://localhost/user/tom/quangle.txt`
- `hadoop fs -copyToLocal quangle.txt quangle.copy.txt`
- `hadoop fs -mkdir books`
- `hadoop fs -ls`
-

# *Hadoop Filesystems*

- Local
- HDFS
- HFTP
- HSFTP
- WebHDFS
- HAR
- KFS (Cloud-Store)
- FTP
- S3 (native)
- S3 (block-based)
- Distributed RAID
- View



## *Filesystem -Local*

- URI scheme -file
- Java implementation - `fs.LocalFileSystem`
- Description -A filesystem for a locally connected disk with client- side checksums. Use `RawLocalFileSystem` for a local filesystem with no checksums.

# *Filesystem - HDFS*

- URI scheme -hdfs
- Java implementation - `hdfs.DistributedFileSystem`
- Description -Hadoop's distributed filesystem. HDFS is designed to work efficiently in conjunction with MapReduce.

## *Filesystem -HFTP*

- URI scheme -hftp
- Java implementation - `hdfs.HftpFileSystem`
- Description -A filesystem providing read-only access to HDFS over HTTP. (Despite its name, HFTP has no connection with FTP.) Often used with `distcp` to copy data between HDFS clusters running different versions.

## *Filesystem - HSFTP*

- URI scheme -hsftp
- Java implementation -hdfs.HsftpFileSystem
- Description -A filesystem providing read-only access to HDFS over HTTPS. (Again, this has no connection with FTP.)

## *Filesystem - WebHDFS*

- URI scheme -webhdfs
- Java implementation -hdfs.web.WebHdfsFile System
- Description -A filesystem providing secure read-write access to HDFS over HTTP. WebHDFS is intended as a replacement for HFTP and HSFTP.

## *Filesystem - HAR*

- URI scheme -har
- Java implementation -fs.HarFileSystem
- Description - A filesystem layered on another filesystem for archiving files. Hadoop Archives are typically used for archiving files in HDFS to reduce the namenode's memory usage.

## *Filesystem -KFS (Cloud-Store)*

- URI scheme -kfs
- Java implementation -fs.kfs.KosmosFileSystem
- Description - CloudStore (formerly Kosmos filesystem) is a distributed filesystem like HDFS or Google's GFS, written in C++

# *Filesystem - FTP*

- URI scheme -ftp
- Java implementation -fs.ftp.FTPFileSystem
- Description - A filesystem backed by an FTP server.



## *Filesystem - S3 (native)*

- URI scheme -s3n
- Java implementation -fs.s3native.NativeS3FileSystem
- Description - A filesystem backed by Amazon S3

## *Filesystem - S3 (block-based)*

- URI scheme -s3
- Java implementation -fs.s3.S3FileSystem
- Description -A filesystem backed by Amazon S3, which stores files in blocks (much like HDFS) to overcome S3's 5 GB file size limit.

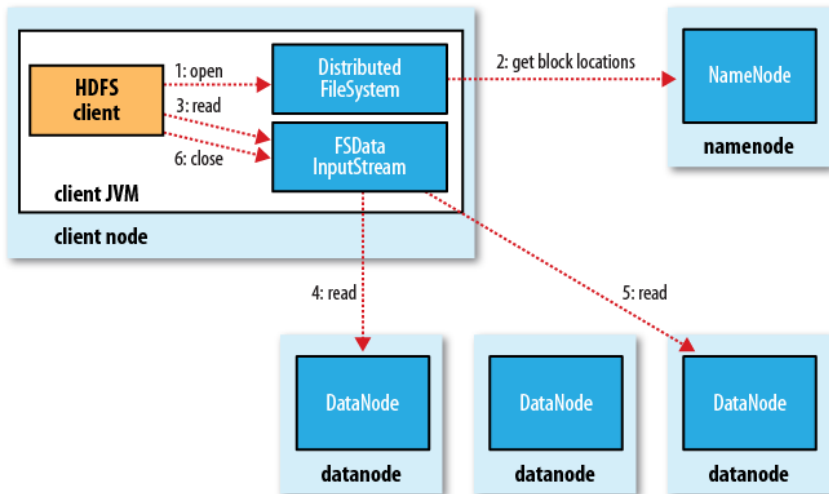
## *Filesystem - Distributed RAID*

- URI scheme -hdfs
- Java implementation -hdfs.DistributedRaidFileSystem
- Description - A “RAID” version of HDFS designed for archival storage. For each file in HDFS, a (smaller) parity file is created, which allows the HDFS replication to be reduced from three to two, which reduces disk usage by 25% to 30% while keeping the probability of data loss the same. Distributed RAID requires that you run a RaidNode daemon on the cluster.

## *Filesystem - View*

- URI scheme -viewfs
- Java implementation -viewfs.ViewFileSystem
- Description -A client-side mount table for other Hadoop filesystems. Commonly used to create mount points for federated namenodes

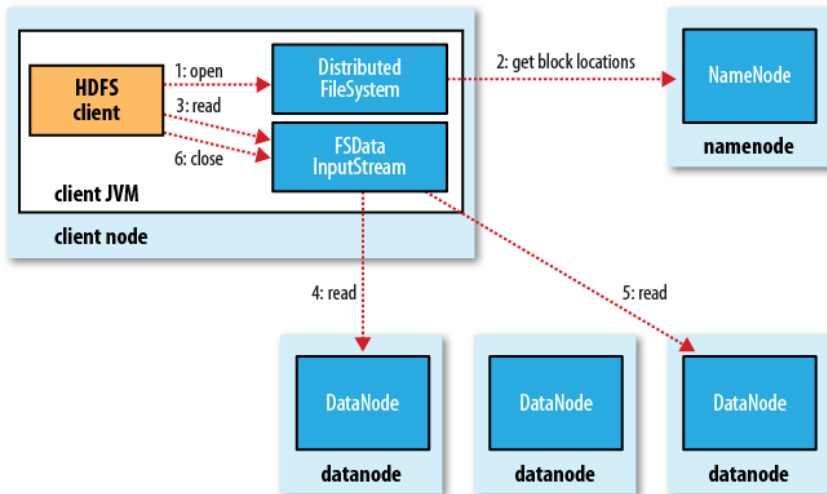
## *Anatomy of a File Read*



## *HDFS-File Read*

- HDFS client sends Read() to File system
- DFS - getblocklocation - Namenode
- client- read()-FSdata input stream
- Read data from Datanodes
- if Datanode down - fetch from next source- Remember the node
- after finishing the set , repeat the steps for moreblocks
- DFSInputStream verify the checksums- if error report to namenode
- Namenode- Inmemory computing

## *Anatomy of a file write*



# *Replica Placement*

- trade-off between reliability and write bandwidth and read bandwidth



## *Parallel Copying with distcp*

Hadoop comes with a useful program called `distcp` for copying data to and from Hadoop filesystems in parallel.

- `hadoop distcp file1 file2`
- `hadoop distcp dir1 dir2`
- `hadoop distcp -update dir1 dir2`

## *Keeping an HDFS Cluster Balanced*

- balancer tool-subsequently even out the block distribution across the cluster.

Table of contents

○

SYLLABUS

○

History

oooooooooooooooooooooooooooo

Anatomy of a File Read

ooooo

